

Computer Assignment (CA) No. 8: Central Limit Theorem

Tyler Berezowsky

March 15, 2015

1 PROBLEM STATEMENT

The goal of this assignment is to demonstrate an application of the Central Limit Theorem. The tasks to be accomplished are:

1. Generate a sum of uniformly distributed mutually independent random variables:

$$S_n = X_1 + X_2 + \cdots + X_n \quad (1.1)$$

Write this as a function in MATLAB with arguments that include the number of random variables (n), the number of total samples generated (N), and the range of the uniform random number generator (e.g., $\text{min}=-1$, $\text{max}=1$).

2. In the main part of your program, write a loop for $n=1, 100$, and call this function for $N=10,000$ with a range of $[-1, 1]$. For each iteration, compute the mean and variance of the output sequence, S_n , and plot the RMS error between a Gaussian fit of this distribution and the actual distribution (I hope you are using your code from a previous homework assignment in a function!).
3. Plot the RMS error as a function of the value of n . Also display the actual distribution and overlay its Gaussian fit for $n=1$, $n=10$ and $n=100$.
4. Consider the following technique for generating a Gaussian distribution from a uniform random number generator:
http://en.wikipedia.org/wiki/Box%E2%80%93Muller_transform.
Generate $N=10,000$ random numbers using this technique, estimate a pdf, and compare the result using the RMS error to (2) for $n=10$, $N=10,000$, $\text{range}=[0, 1]$. Time the code in both cases using MATLABs built-in timing tools. Which technique gives the better fit? Which technique is faster? How low can you set n to get comparable performance in both time and RMS error?

Discuss how (1)-(3) demonstrates the Central Limit Theorem.

2 APPROACH AND RESULTS

2.1 TASKS 1 TO 3

A function `uniformsum(n, N, a, b)` was constructed which generates n random variables with a length of N and limits equal to $[b, a]$. This function was called with $n = [1, 100]$ and $N = 10000$, and for each value of n the PDF calculated from the resulting vector S_n . The normal distribution was also fitted per n . The RMS error, defined below, was calculated for each value of n and plotted as a function of n where f_1 is the fitted normal distribution and f_2 is the PDF of S_n . This can be seen in figure 2.1.

$$RMS_{error} = \sqrt{MSE} = \sqrt{E[(f_1 - f_2)^2]} \quad (2.1)$$




Figure 2.1: \sqrt{MSE} as a function of n

Interestingly, the \sqrt{MSE} appears to converge between 0.06 and 0.04 after only ten iterations which infers diminishing returns after $n = 10$. Plots of the PDT_{S_n} are displayed in figure 2.2 with $n = \{1, 10, 100\}$. As previously stated, n is equal to number of uniformly distributed random variables of length N being summed to generate S_n .

When $n = 1$, S_n represents a single random variable, and is thus a uniform distribution. This is illustrated in the top subplot. When $n = 10$, S_n represents the sum of 10 random variables. As the middle subplot illustrates, the distribution of S_n now approaches a normal distribution. When $n = 100$, S_n represents the sum of 100 random variables. The PDF of S_n is displayed in the bottom subplot, and as predicted by error plot in figure 2.1, the distribution appears identical in approximation to the fitted normal distribution of S_n .




Figure 2.2: Histogram of S_n and fitted normal distribution

2.2 TASK 4

The Box-Muller method was utilized to generate a normally distributed random variable and timed against the calculation for S_n with $n = [1, 10]$. The Box-Muller method is described below where Z_0 and Z_1 are normally distributed random variables, and U_1 and U_2 are independent uniformly distributed random variables:

$$Z_0 = R \cos(\theta) = \sqrt{-2 \ln U_1} \cos(2\pi U_2) \quad (2.2)$$

$$Z_1 = R \sin(\theta) = \sqrt{-2 \ln U_1} \sin(2\pi U_2) \quad (2.3)$$

The Box-Muller method and calculation for S_n was trialed 10 times for each value of n . The result was averaged per n . The RMS error between the fitted normal distribution, S_n , and the Box-Muller method was also determined. The results are illustrated in figure 2.3.

The top subplot illustrates the time for each function given n . As expected the time to complete the calculation for S_n increases with n . The Box-Muller method is constant as n is irrelevant for the calculation. The functions intersect when $n \approx 2$. The bottom subplot illustrates the error between the normal distribution and the PDF of the Box-Muller method and S_n . As previously described, the RMS between S_n and the normal decreases with n . The Box-Muller method is constant because n is irrelevant in the calculation. The error of the Box-Muller method is consistently lower than that of S_n as seen for $n = [1, 10]$. Figure 2.1 displays that the error converges to 0.06 and 0.04 as n increases beyond 10, therefore it can be expected that the Box-Muller will be more accurate despite increasing values of n . Further trials would have to be attempted to confirm the previous assumption, but if error for S_n does decrease with n , the time required would be magnitudes larger than that of the Box-Muller method.

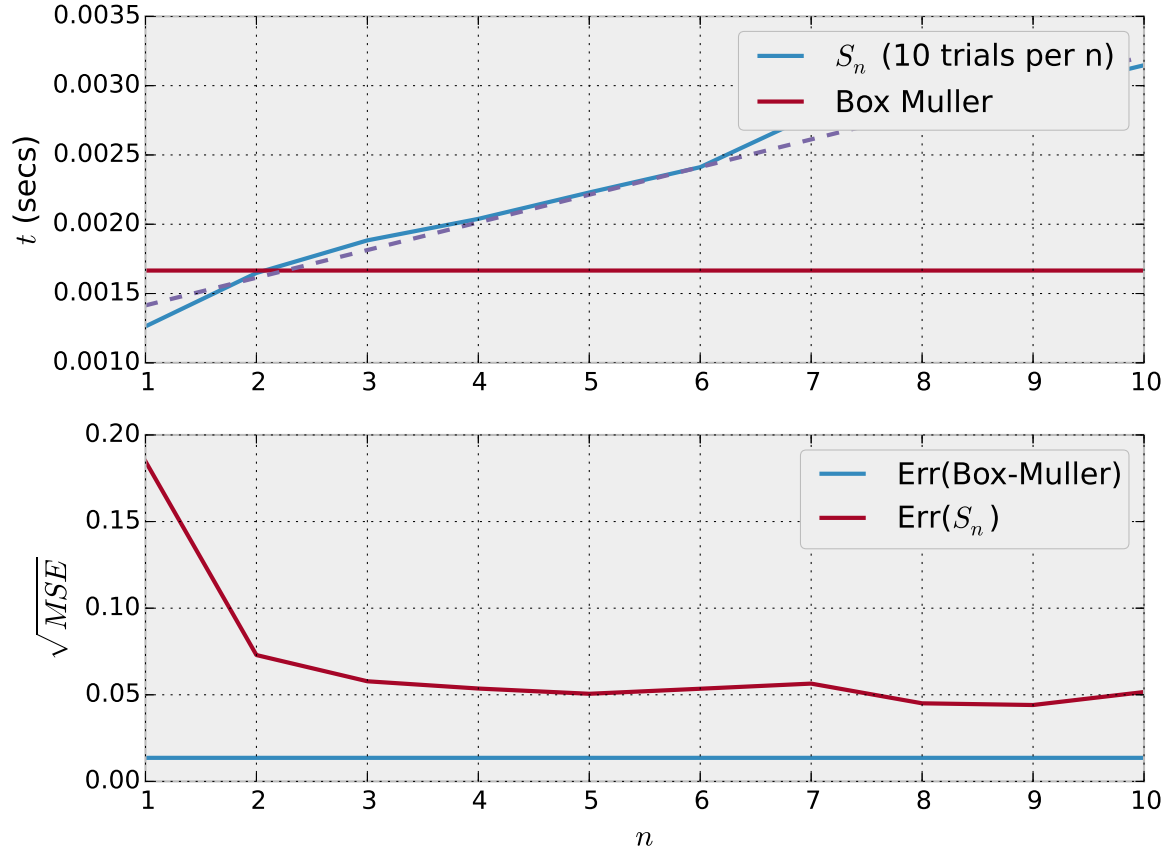


Figure 2.3: Comparison of calculation for S_n and Box-Muller method for the generation of a random normally distributed vector.

3 SOURCE CODE

```

1 import numpy as np
2 import scipy.stats as stats
3 from numpy.random import rand
4 import matplotlib.pyplot as plt
5 import time
6 plt.style.use('bmh')
7
8
9 def uniformSum(n, N, a, b):
10     X_n = (b - a)*rand(N, n) + a # generate uniform RVs between b and a
11     S_n = np.sum(X_n, axis=1) # sum
12     S_n = S_n/np.max(abs(S_n))
13     return S_n
14
15
16 def normal(x, mu, variance):
17     pdf = (1./np.sqrt(2*np.pi*variance))*np.exp(-(x-mu)**2/(2*variance))
18     return pdf
19
20

```

```

21 def findcenterbins(bins):
22     centerBins = np.zeros(len(bins) - 1)
23     for index in range(0, len(bins) - 1):
24         centerBins[index] = np.mean([bins[index + 1], bins[index]])
25     return centerBins
26
27
28 def MSE(f1, f2):
29     mse = np.mean((f1-f2)**2)
30     return mse
31
32
33 # PARAMETERS
34 N = 10000
35 a = -1
36 b = 1
37 BINS = 100
38 NSUMS = 100
39
40 # MEMORY ALLOCATION
41 S_n = np.zeros(N)
42 error = np.zeros(NSUMS)
43 for n in range(1, NSUMS+1):
44     S_n = uniformSum(n, N, a, b) # generate sum
45     sVar = np.var(S_n) # compute variance
46     sMu = np.mean(S_n) # compute mean
47     sPDF, bins = np.histogram(S_n, bins=BINS, normed=True) # compute PDF
48     centerbins = findcenterbins(bins) # compute the center of bins
49     sNormal = normal(centerbins, sMu, sVar) # compute norm dis via metrics
50
51     error[n-1] = np.sqrt(MSE(sPDF, sNormal))
52
53 # PLOTS
54 n = range(1, NSUMS+1)
55 eplt = plt.figure(1)
56 ax = eplt.add_subplot(111)
57 ax.plot(n, error, label=r'$\sqrt{MSE}$')
58 ax.legend()
59 ax.set_xlabel(r'$n$')
60 plt.show()
61
62 # plot n = {1, 10, 100}
63 ns = [1, 10, 100]
64 hplt, axs = plt.subplots(3, 1, sharex=True, sharey=True)
65 for n, ax in zip(ns, axs):
66     S_n = uniformSum(n, N, a, b)
67     sVar = np.var(S_n) # compute variance
68     sMu = np.mean(S_n) # compute mean
69     centerbins = findcenterbins(bins) # compute the center of bins
70     sNormal = normal(centerbins, sMu, sVar) # compute norm dis via metrics
71

```

```

72     ax.hist(S_n, bins=BINS, normed=True, alpha=0.5, label=r'$PMF_{S_{\{d\}}}$' % n)
73     ax.plot(centerbins, sNormal, label=r'$N(S_{\{d\}})$' % n)
74     ax.legend()
75
76
77 # BOX-Muller Transform
78 def boxmuller(N):
79     a = 0; b = 1; n = 1
80     U1 = (b - a)*rand(N, n) + a
81     U2 = (b - a)*rand(N, n) + a # generate uniform RVs between b and a
82     R = np.sqrt(-2*np.log(U1))
83     theta = 2*np.pi*U2
84     Z1 = R*np.cos(theta)
85     return Z1
86
87 # TIME TEST
88
89 N = 10000
90 ns = range(1, 11)
91 trials = range(0, 10)
92 start = zeros(10)
93 end = zeros(10)
94 tstart = zeros(10)
95 tend = zeros(10)
96
97 for n in ns:
98     for trial in trials:
99         tstart[trial] = time.time()
100         uniformSum(n, N, 0, 1)
101         tend[trial] = time.time()
102         start[n-1] = np.mean(tstart)
103         end[n-1] = np.mean(tend)
104 utime = end - start
105
106 S_n = []
107 error2 = []
108 for trial in trials:
109     start[trial] = time.time()
110     S_n = boxmuller(N)
111     end[trial] = time.time()
112 print S_n
113 sVar = np.var(S_n) # compute variance
114 sMu = np.mean(S_n) # compute mean
115 sPDF, bins = np.histogram(S_n, bins=BINS, normed=True) # compute PDF
116 centerbins = findcenterbins(bins) # compute the center of bins
117 sNormal = normal(centerbins, sMu, sVar) # compute norm dis via metrics
118 error2 = np.sqrt(MSE(sPDF, sNormal))
119
120 btime = np.mean(end - start)
121
122 tplot = plt.figure(3)

```

```

123 ax1 = tplot.add_subplot(211)
124 ax2 = tplot.add_subplot(212)
125 ax1.plot(ns, utime, label='$S_n$ (10 trials per n)')
126 ax1.plot(ns, [btime]*10, label='Box Muller')
127 ax2.plot(ns, [error2]*len(ns), label=r'Err (Box-Muller)')
128 n = np.arange(1, 11)
129 ax2.plot(n, error[0:len(n)], label='Err ($S_n$)')
130 z = np.polyfit(ns, utime, 1)
131 p = np.poly1d(z)
132 ax1.plot(ns, p(ns), '--')
133 ax1.legend()
134 ax2.set_xlabel(r'$n$')
135 ax1.set_ylabel(r'$t$ (secs)')
136 ax2.set_ylabel(r'$\sqrt{\text{MSE}}$')
137
138 ax2.legend()
139 tplot.tight_layout()

```

4 CONCLUSIONS

Due to the Central Limit Theorem, as n is increased the sum of n uniformly distributed random variables approaches a vector which is normally distributed. Central Limit Theorem: The arithmetic mean of a sufficient number of independent random variables will be *approximately* normally distributed given the random variables are identically distributed.¹

This theorem can be utilized to generate a random normally distributed variable, but as displayed takes considerable effort in comparison to other methods such as the Box-Muller method. Further exploration is required to understand why the RMS error appears to converge as n increases. The proof of the Central Limit Theorem states given any random variable Y with $E(Y) = 0$ and $var(Y) = 1$ summed n times as $\lim_{n \rightarrow \infty}$ the approximation approaches $e^{-t/2}$ which would be equal to $N(0, 1)$. Therefore, error should decrease as n increases.

¹http://en.wikipedia.org/wiki/Central_limit_theorem